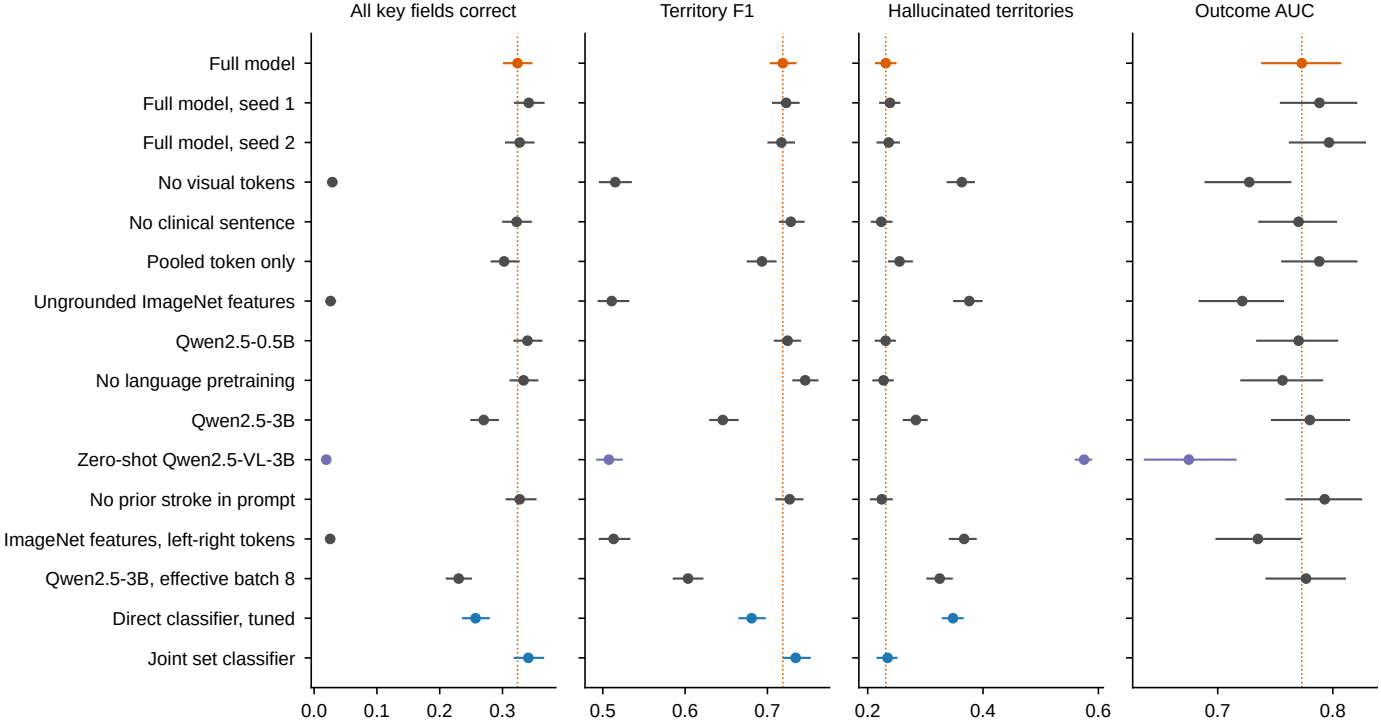
All key fields correct
Territory F1
Hallucinated territories
Outcome AUC
Full model
Full model, seed 1
Full model, seed 2
No visual tokens
No clinical sentence
Pooled token only
Ungrounded ImageNet features
Qwen2.5-0.5B
No language pretraining
Qwen2.5-3B
Zero-shot Qwen2.5-VL-3B
No prior stroke in prompt
ImageNet features, left-right tokens
Qwen2.5-3B, effective batch 8
Direct classifier, tuned
Joint set classifier
0.0
0.1
0.2
0.3
0.5
0.6
0.7
0.2
0.4
0.6
0.7
0.8